

# AIFA: Stochastic Planning MDP

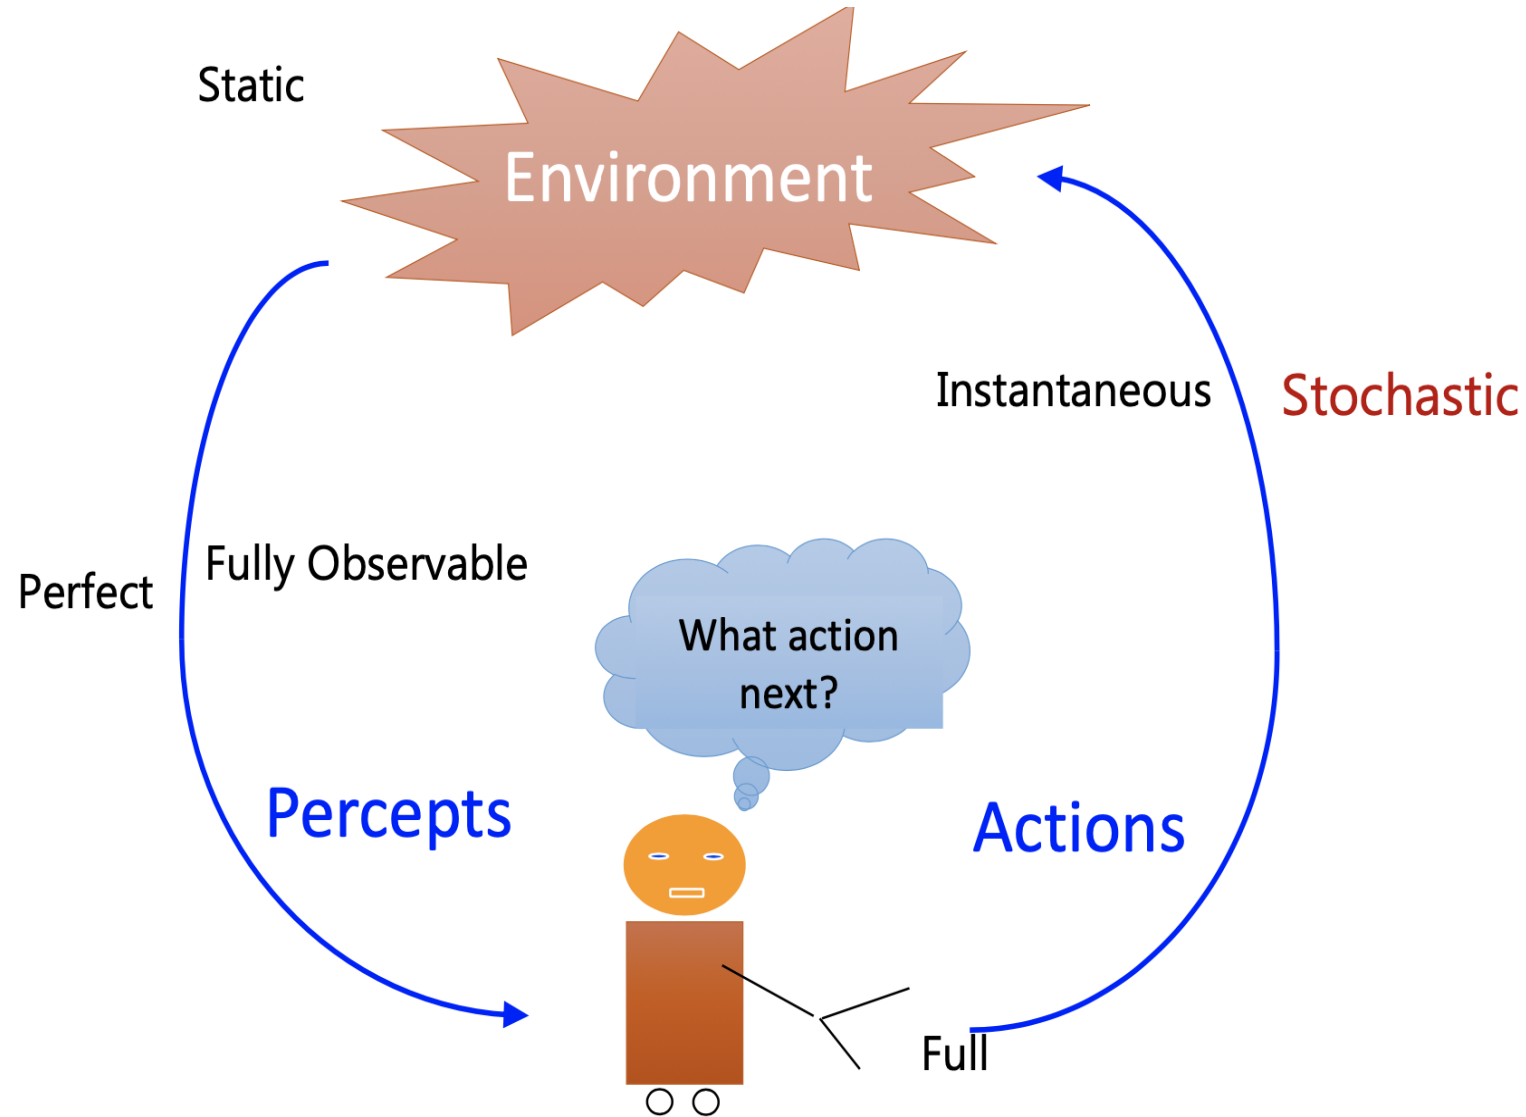
01/04/2025

**Koustav Rudra**


# Markov Decision Processes

- Value Iteration
- Policy Iteration

# Planning under Uncertainty



# Uncertainty



Randomness:



- could be caused by limitations of the sensors and actuators of the robot
- could be caused by market forces or nature

...

- **Robotics**: decide where to move, but actuators can fail, hit unseen obstacles, etc.
- **Resource allocation**: decide what to produce, don't know the customer demand for various products
- **Agriculture**: decide what to plant; don't know weather and thus crop yield


# Example: Grid World

- **Noisy movement: actions do not always go as planned**
  - 80% of the time, the action North takes the agent North (if there is no wall there)
  - 10% of the time, North takes the agent West; 10% East
  - If there is a wall in the direction the agent would have been taken, the agent stays put
- **The agent receives rewards**
  - Small “living” reward each step
  - Big rewards come at the end
- **Goal: maximize sum of rewards**




# Grid World Actions

## Deterministic Grid World




## Stochastic Grid World




# Markov Decision Process

- An MDP is defined by:
  - A **set of states**  $s \in S$
  - A **set of actions**  $a \in A$
  - A **transition function**  $T(s, a, s')$ 
    - Probability that  $a$  from  $s$  leads to  $s'$ , i.e.,  $P(s'|s, a)$
    - Also called the model or the dynamics
  - A **reward function**  $R(s, a, s')$ 
    - Sometimes just  $R(s)$  or  $R(s')$
  - A **start state**
  - Maybe a **terminal state**



# Markov Decision Process

- An MDP is defined by:
  - A **set of states**  $s \in S$
  - A **set of actions**  $a \in A$
  - A **transition function**  $T(s, a, s')$ 
    - Probability that  $a$  from  $s$  leads to  $s'$ , i.e.,  $P(s'|s, a)$
    - Also called the model or the dynamics
  - A **reward function**  $R(s, a, s')$ 
    - Sometimes just  $R(s)$  or  $R(s')$
  - A **start state**
  - Maybe a **terminal state**




Goal: Learn to choose actions that maximize reward




# Solution to MDP: Policies

- For MDPs, we want an optimal **policy**  $\pi^*: S \rightarrow A$
- A policy  $\pi$  gives an action for each state
- An optimal policy is one that maximizes expected utility if followed

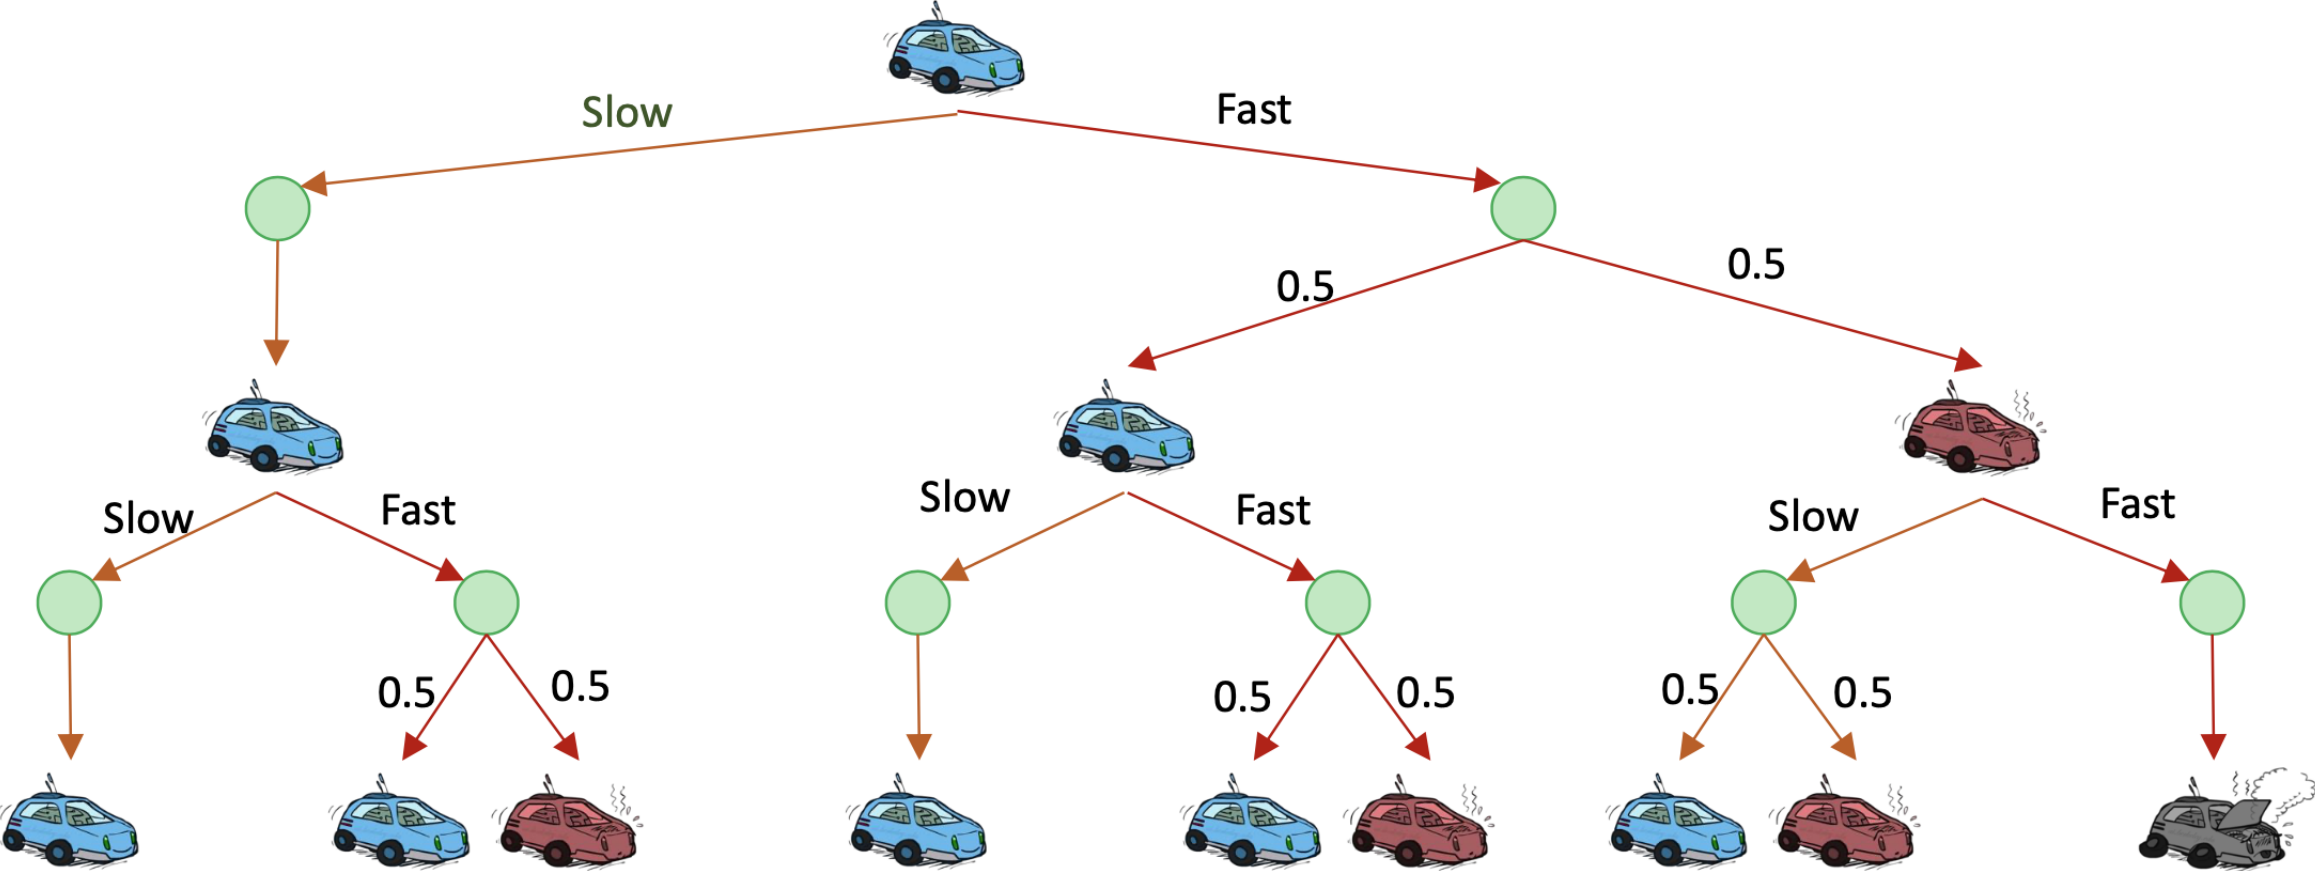


# Example: Racing

- A robot car wants to travel far, quickly
- Three states: **Cool**, **Warm**, Overheated
- Two actions: **Slow**, **Fast**
- Going faster gets double reward




# Racing car search tree



# MDP search trees

Each MDP state projects an expectimax-like search tree



# Utilities

## Two ways to define utilities


- Additive utility:  $U([r_0, r_1, r_2, \dots]) = r_0 + r_1 + r_2 + \dots$
- Discounted utility:  $U([r_0, r_1, r_2, \dots]) = r_0 + \gamma r_1 + \gamma^2 r_2 \dots$   
 $0 \leq \gamma \leq 1$ : discount factor

# What is a solution?

- A policy  $\pi$  is a mapping from each state  $s \in \text{States}$  to an action  $a \in \text{Actions}(s)$
- Evaluating a policy
  - Following a policy yields a random path
  - The **utility** of a policy is the (discounted) sum of the rewards on the path (this is a random variable).
  - The **value** of a policy at a state is the expected utility.


# Policy Evaluation

- Definition: value of a policy
  - Let  $V_\pi(s)$  be the expected utility received by following policy  $\pi$  from state  $s$
- Definition: Q-value of a policy
  - Let  $Q_\pi(s, a)$  be the expected utility of taking action  $a$  from state  $s$ , and then following policy  $\pi$



# Policy Evaluation

- **Plan:** define recurrences relating value and Q-value



$$V_\pi(s, a) = \begin{cases} 0 & \text{if } \text{isEnd}(s) \\ Q_\pi(s, \pi(s)) & \text{otherwise} \end{cases}$$

$$Q_\pi(s, a) = \sum_{s'} P(s'|s, a) [R(s, a, s') + \gamma V_\pi(s')]$$




# Policy Evaluation

- Iterative algorithm: Start with arbitrary policy values and repeatedly apply recurrences to converge to true values
- Initialize  $V_{\pi}^0(s) \leftarrow 0$  for all states  $s$
- For iteration  $t = 1, \dots, t_{PE}$ 
  - For each state  $s$ 
    - $V_{\pi}^{(t)}(s) \leftarrow \sum_{s'} T(s, \pi(s), s') [R(s, \pi(s), s') + \gamma V_{\pi}^{(t-1)}(s')]$
- How many iterations  $t_{PE}$ ?
  - Repeat until values do not change much
  - $\max_s |V_{\pi}^{(t)}(s) - V_{\pi}^{(t-1)}(s)| \leq \epsilon$

# Optimal Quantities


- The value (utility) of a state  $s$ :
  - $V^*(s) =$   
*expected utility starting in  $s$  and acting optimally*
- The value (utility) of a q-state  $(s,a)$ :
  - $Q^*(s,a) =$   
*expected utility starting out having taken action  $a$  from state  $s$  and thereafter acting optimally*
- The optimal policy:
  - $\pi^*(s) =$  optimal action from state  $s$



# Value Function

- Value function for a policy  $\pi: S \rightarrow A$ 
  - $V^\pi(s) = R(s) + \gamma \sum_{s' \in S} T(s, \pi(s), s') V^\pi(s')$
- Optimum value function
  - $V^*(s) = \max_{\pi} V^\pi(s)$
- $Q^\pi(s, a)$ : the expected utility of taking action  $a$  from state  $s$ , and then following policy  $\pi$ .

# Optimal Values and Q-values


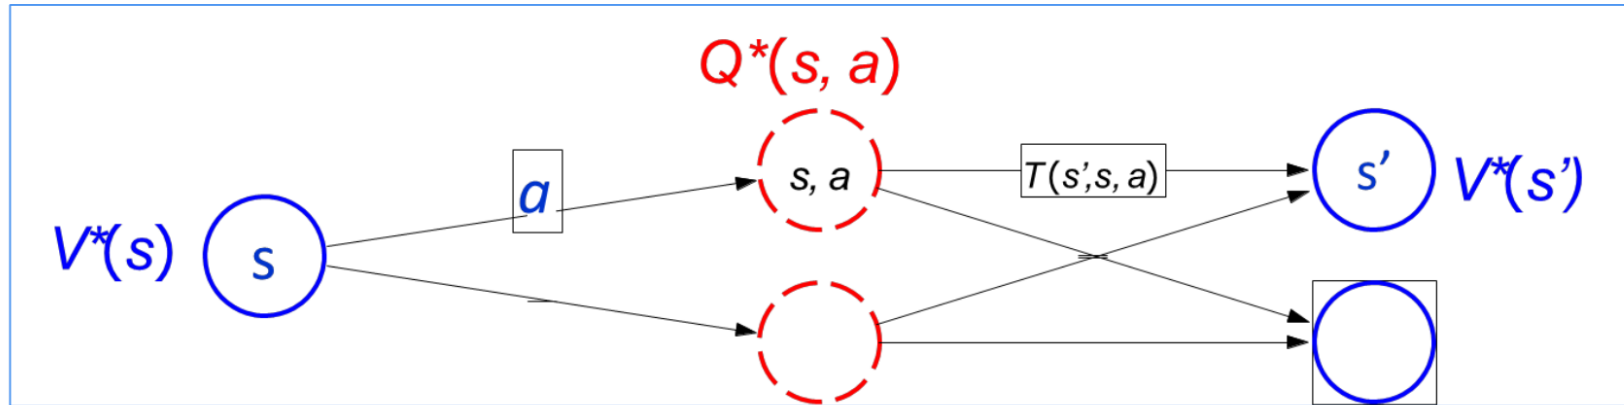


$$Q^*(s, a) = \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V^*(s')]$$

$$V^*(s) = \max_a Q^*(s, a)$$

# The Bellman Equations

- Definition of “optimal utility” leads to a simple one-step lookahead relationship amongst optimal utility values:



$$V^*(s) = \max_a Q^*(s, a)$$

$$Q^*(s, a) = \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V^*(s')]$$

$$V^*(s) = \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V^*(s')]$$

Thank You